



How to Cook with C++

Preparation

- Have you ever wanted to do something **useful** with C++?
- Do you feel like you aren't really using the **power** of you machine?
- Do you like eggs?

What we will Need



A GOOD GRILL



A GOOD GRILL
PROGRAM



A FEW EGGS

Step 1: Prepare the Grill




- What constitutes a good grill program?
 - Tons of Loads/Stores ? NO
 - Multithreading ? Yes
 - Vectorization ? Yes
 - High Flops? Yes

Step 2: Write a Grill Program

- Ok time for Grill_v1.cpp

Grill_v1.cpp

```
int main() {  
    while (true) {  
        double i = 0.;  
        i += 1.36236236236;  
        i *= 236236.23623699102;  
    }  
}
```




g++ -O3 grill_v1.cpp

We are not cooking with heat at all...

Grill_v1_parallel.cpp

```
int main() {  
    while(true) {  
#pragma omp parallel for  
        for (int i = 0; i < 1000; ++i) {  
            i += 1.362362326;  
            i *= 2336236.68236;  
        }  
    }  
}
```



Utilizing all CPUs

Temp is better but still low

This Grill Program Kinda Sucks



Temperature to cook egg $\sim 65^{\circ}\text{C}$;
need higher temps to account
for heat loss



It's Not achieving **max flops**...




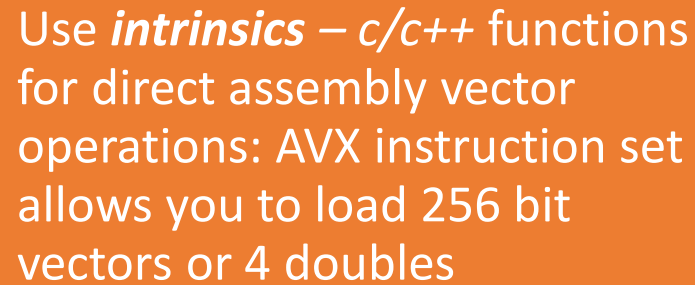
Let's do some **math**...

My Grill


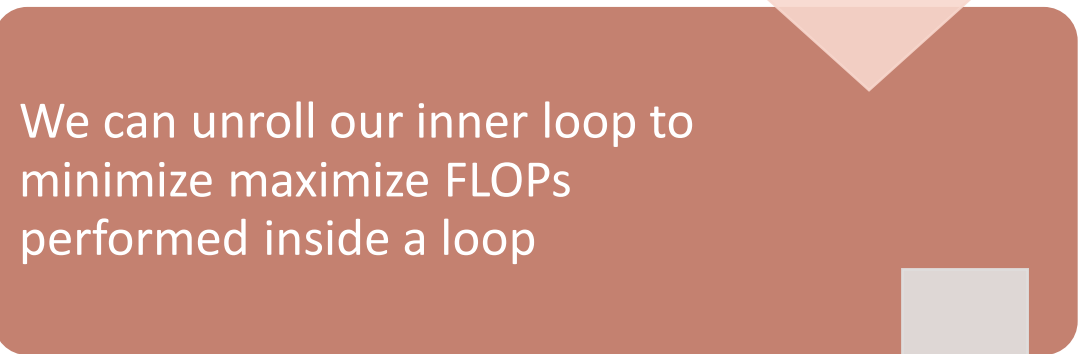
- 2.6 Ghz intel processor, support for AVX 256 vector operations
- Max flops is $2.6 * 8 \text{ flops} = 20.8 \text{ Gflops}$
- Max total flops over **6 cores**: $6 * 2.6 * 8 = 124.8 \text{ Gflops}$
- Let's give it a go...

How to Cook with Gasline

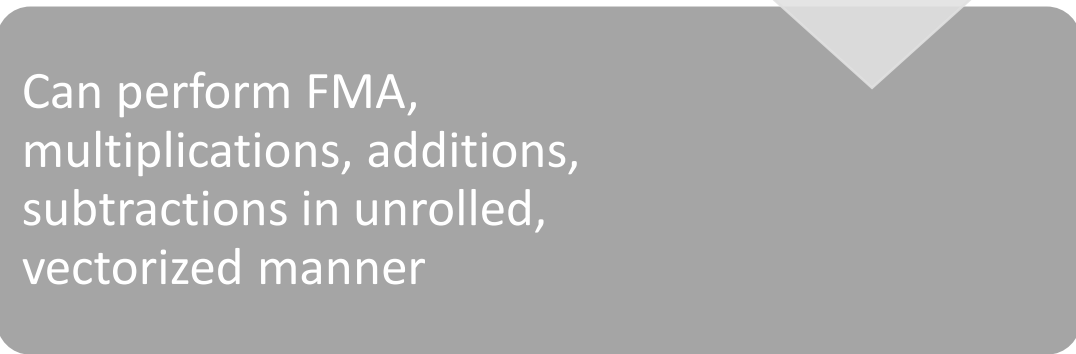
Use *intrinsics* – c/c++ functions for direct assembly vector operations: AVX instruction set allows you to load 256 bit vectors or 4 doubles



We can unroll our inner loop to minimize maximize FLOPs performed inside a loop



Can perform FMA, multiplications, additions, subtractions in unrolled, vectorized manner



Red Hot Grill

d1	d2	d3	d4
----	----	----	----

Vector Intrinsics

`__mm256d vector = __mm256_set1_pd(5)` – Broadcast a double to a 4 double wide vector

`__mm256_fmadd_pd(__mm256d v1, __mm256d v2, __mm256d v3)` – Fused Multiply Add

`__mm256_add_pd(__mm256d v1, __mm256d v2)` – Vectorized Add

`__mm256_mul_pd(__mm256d v1, __mm256d v2)` – Vectorized Multiply

```
_mm256_add_pd (v1, v2);
```

```
_mm256_mul_pd (v1, v2)
```


```
_mm256_mul_pd (v1, v2)
```

```
_mm256_add_pd (v1, v2);
```

```
_mm256_mul_pd (v1, v2)
```

....

UNROLLED





Voila